

TLX: Minimally Invasive Paths to Performance Portability

Hongtao Yu (Meta)



Outline

- Motivation
- Overview of TLX
- Building blocks
- Kernel performance study
- What's Next
- Key takeaways



Extending Triton's Power

- Triton is great
 - Exceptional productivity through tile-level abstraction and automatic thread mapping
 - Hardware-agnostic design — write once, run anywhere
 - Delivers effortless high performance on CUDA cores
- Triton can be greater if
 - Attainable path to ultimate performance on tensor cores
 - Stable and predictable performance in production
 - Fast path to new hardware features amid short GPU lifecycles



TLX : Triton Low-level Language Extension

- Extends Triton DSL with low-level hardware-specific primitives
- Exposes fine-grained control over hardware features for performance
- (+) Abstracts away unnecessary hardware details (e.g., layout encoding)
- (+) Seamless integration with Triton DSL and compiler
- (+) Preserves Triton's composable semantics
- (+) Enables incremental development and performance tuning



Example: Dot Compression++

Computes $X @ (X^T @ Y)$

```
@triton.jit
def dcpp(x_desc, y_desc, o_desc, BLOCK_K: tl.constexpr,
):
    pid = tl.program_id(0)

    # compute  $X^T @ Y$ 
    for k in range(0, tl.cdiv(K, BLOCK_K)):
        x = x_desc.load([k + pid * K, 0])
        y = y_desc.load([k + pid * K, 0])
        acc = tl.dot(x.T, y, acc)

    acc = acc.to(tl.float16)

    # compute  $X @ (X^T @ Y)$ 
    for k in range(0, tl.cdiv(K, BLOCK_K)):
        x = x_desc.load([k + pid * K, 0])
        acc = tl.dot(x, acc)
    o_desc.store([k + pid * K, 0], out.to(tl.float16))
```



Example: Dot Compression++

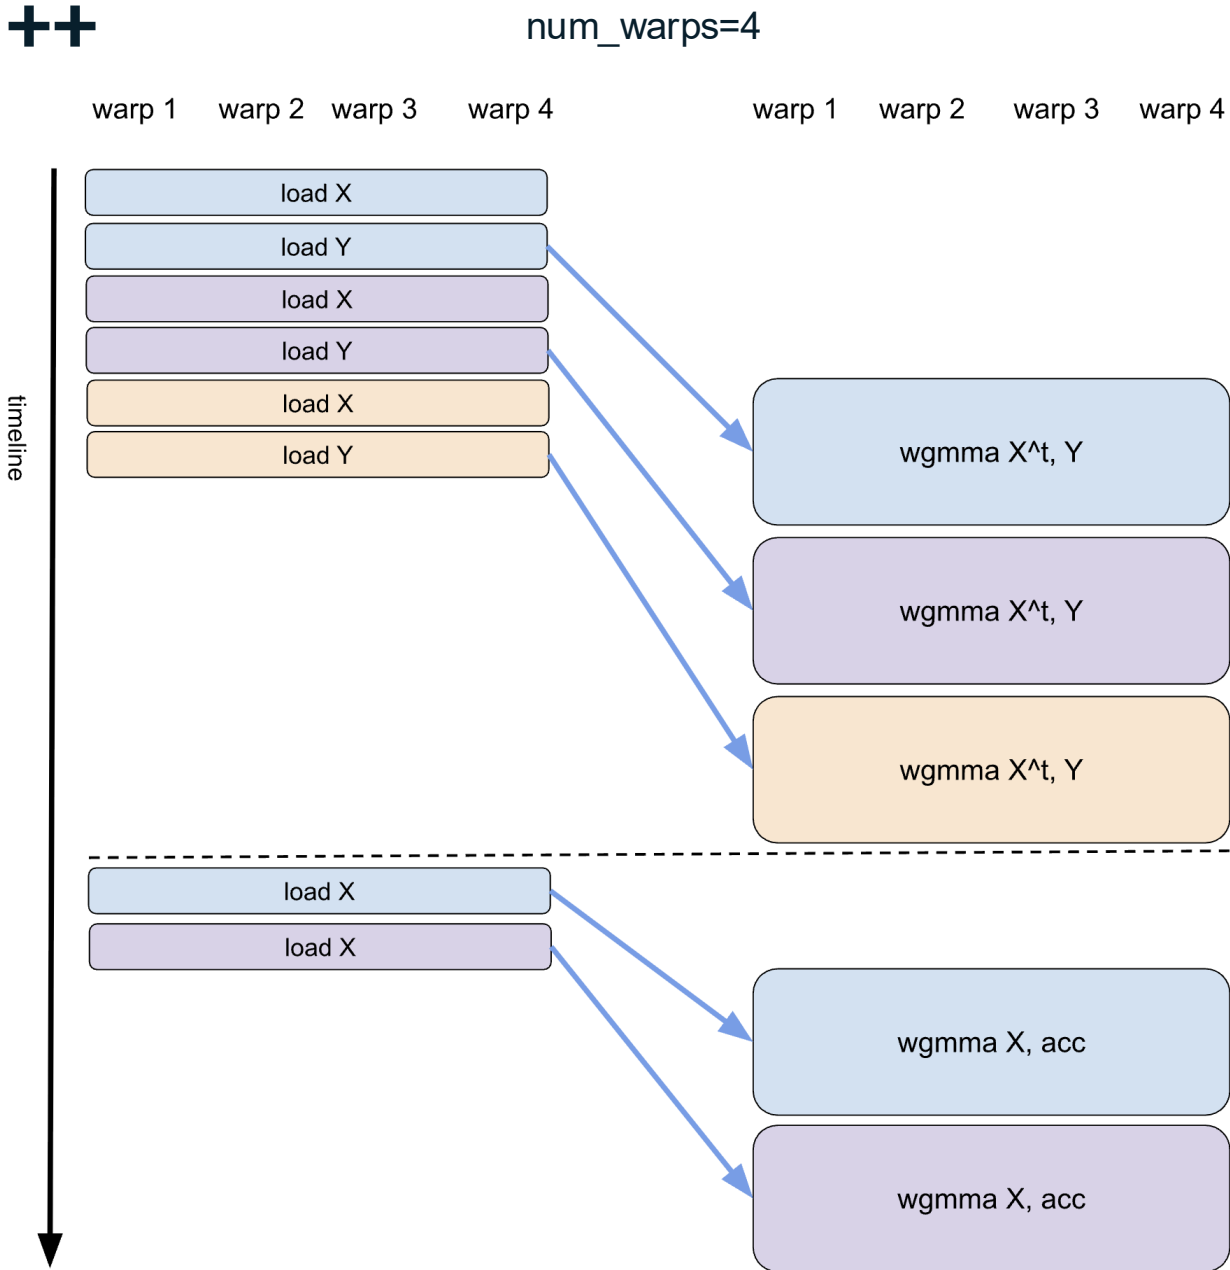
Computes $X @ (X^T @ Y)$

```
@triton.jit
def dcpp(x_desc, y_desc, o_desc, BLOCK_K: tl.constexpr,
):
    pid = tl.program_id(0)

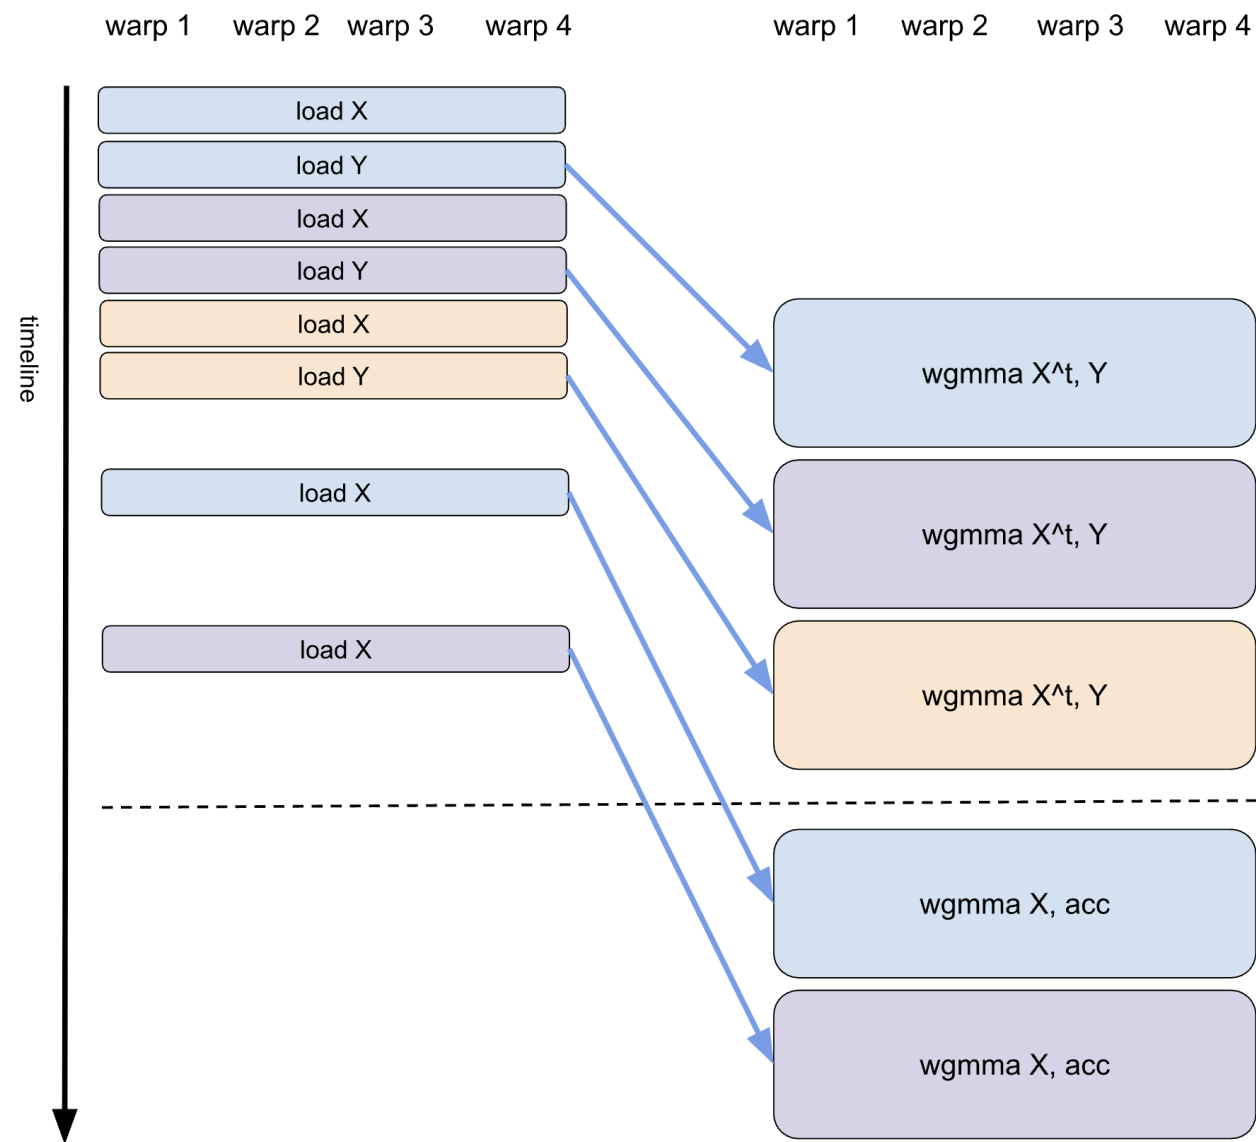
    # compute X^T @ Y
    for k in range(0, tl.cdiv(K, BLOCK_K)):
        x = x_desc.load([k + pid * K, 0])
        y = y_desc.load([k + pid * K, 0])
        acc = tl.dot(x.T, y, acc)

    acc = acc.to(tl.float16)

    # compute X @ (X^T @ Y)
    for k in range(0, tl.cdiv(K, BLOCK_K)):
        x = x_desc.load([k + pid * K, 0])
        acc = tl.dot(x, acc)
        o_desc.store([k + pid * K, 0], out.to(tl.float16))
```



Continuous Pipelining



**20%
speedup**



Example: Dot Compression++

Computes $X @ (X^T @ Y)$

```
@triton.jit
def dcpp(x_desc, y_desc, o_desc, BLOCK_K: tl.constexpr,
):
    pid = tl.program_id(0)

    # compute  $X^T @ Y$ 
    for k in range(0, tl.cdiv(K, BLOCK_K)):
        x = x_desc.load([k + pid * K, 0])
        y = y_desc.load([k + pid * K, 0])
        acc = tl.dot(x, y, acc)

    acc = acc.to(tl.float16)

    # compute  $X @ (X^T @ Y)$ 
    for k in range(0, tl.cdiv(K, BLOCK_K)):
        x = x_desc.load([k + pid * K, 0])
        acc = tl.dot(x, acc)
        o_desc.store([k + pid * K, 0], out.to(tl.float16))
```

```
@triton.jit
def dcpp_tl(x_desc, y_desc, o_desc ..., NB: tl.constexpr):
    pid = tl.program_id(0)
    xt = tl.local_alloc((BM, BK), tl.dtype_of(x_desc), NB)
    y = tl.local_alloc((BK, BN), tl.dtype_of(y_desc), NB)
    x = tl.local_alloc((BM, BK), tl.dtype_of(x_desc), NB, reuse=x)
    bar = tl.alloc_barriers(NB)

    # prefetch  $X^T$  and  $Y$ 
    for k in tl.static_range(0, NB - 1):
        _load_x_y(x_desc, y_desc, xt, y, bar, k, k + pid * K)

    k_iters = tl.cdiv(K, BK)
    for k in range(0, k_iters - NB + 1):
        acc = _dot_xt_y(x, y, acc, bar, k)
        # prefetch  $X^T$  and  $Y$ 
        _load_x_y(x_desc, y_desc, xt, y, bar, k, k + pid * K)

    for k in tl.static_range(k_iters - NB + 1, k_iters):
        acc = _dot_xt_y(x, y, acc, bar, k)
        # prefetch  $X$ 
        _load_x(x_desc, x, bar, k + NB - 1, k - k_iters + NB - 1 + pid * K)

    acc = tl.async_dot_wait(0, acc)
    acc = acc.to(tl.float16)
    for k in range(0, tl.cdiv(K, BK)):
        tlx.barrier_wait(bar[k], (k + NB - 1) // NB)
        out = tlx.async_dot(x[k], y[k], acc)
        _load_x(x_desc, x, bar, k + NB - 1, k - k_iters + NB - 1 + pid * K)
        out = tlx.async_dot_wait(0, out)
        o_desc.store([k + pid * K, 0], out.to(tl.float16))
```

```
@triton.jit
def _load_x_y(x_desc, y_desc, x, y, bar, idx, offset)
    buf = idx % NB
    tlx.barrier_expect_bytes(bar[buf], (BM + BN) * BK)
    tlx.async_descriptor_load(x_desc, x[buf], [offset, 0], bar[buf])
    tlx.async_descriptor_load(y_desc, y[buf], [offset, 0], bar[buf])
```


```
@triton.jit
def _dot_xt_y(x, y, acc, bar, idx)
    buf = idx % NB
    tlx.barrier_wait(bar[buf], phase = buf // NB)
    acc = tlx.async_dot(tlx.local_trans(x[buf]), y[buf], acc)
    return tlx.async_dot_wait(1, acc)
```

```
@triton.jit
def _load_x(x_desc, x, bar, idx, offset)
    buf = idx % NB
    tlx.barrier_expect_bytes(bar[buf], 2 * BM * BK)
    tlx.async_descriptor_load(x_desc, x[buf], [offset, 0], bar[buf])
```



TLX Operations for NVGPU

- Local memory accesses
- Asynchronous global memory accesses
- Asynchronous tensor core operations
- Asynchronous task operation for warp specialization
- Cross-warp group synchronization



Automatic Layout Encoding Assignment

- For register tensors (**tl.tensor**)
 - Non-coalesced placeholder layout introduced during TTIR -> TTGIR
 - Optimized by the memory coalescing pass; or
 - Inferred based on the MMA operations;
 - Reconciled by the removal of layout conversion pass
- For local-memory buffers (**tlx.buffered_tensor**)
 - Non-swizzled placeholder layout introduced during DSL -> TTIR
 - Required by consumers (MMA operations, ...)
 - Propagated and reconciled by layout propagation pass, priority-aware



Outline

- Motivation
- Overview of TLX
- **Building blocks**
- Kernel performance study
- What's Next
- Key takeaways

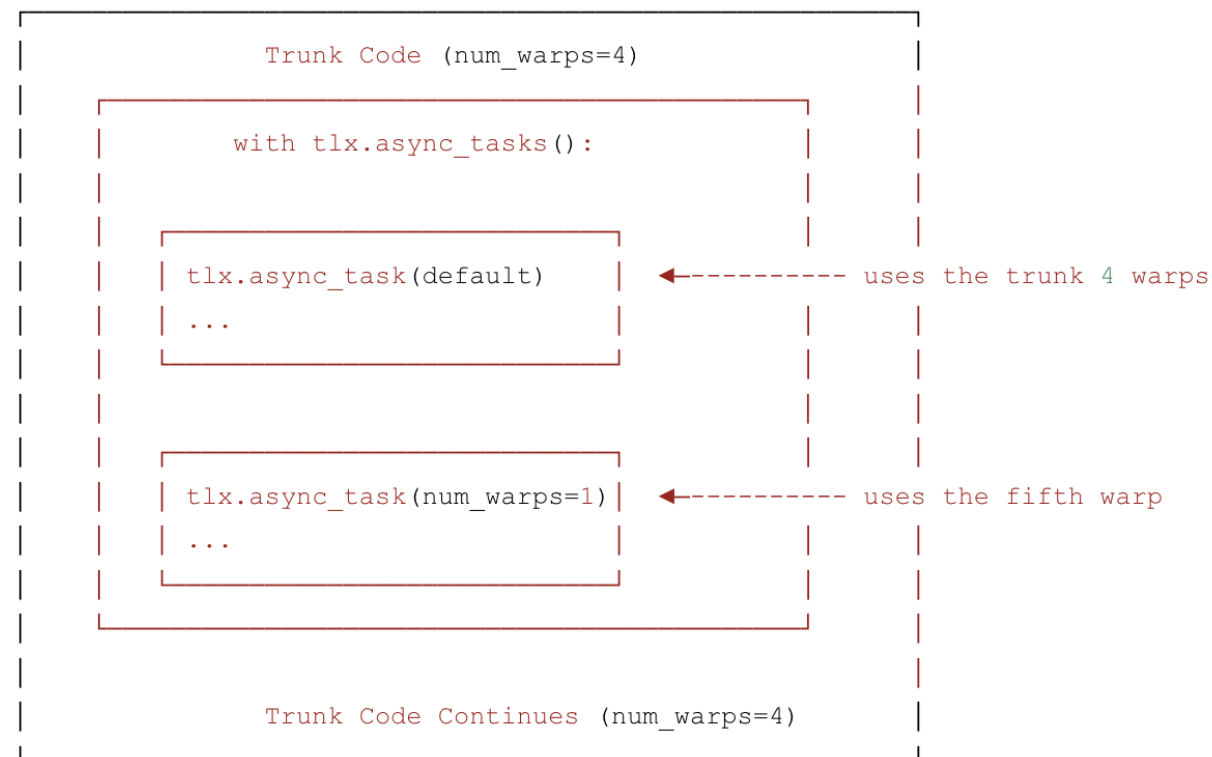


Warp Specialization

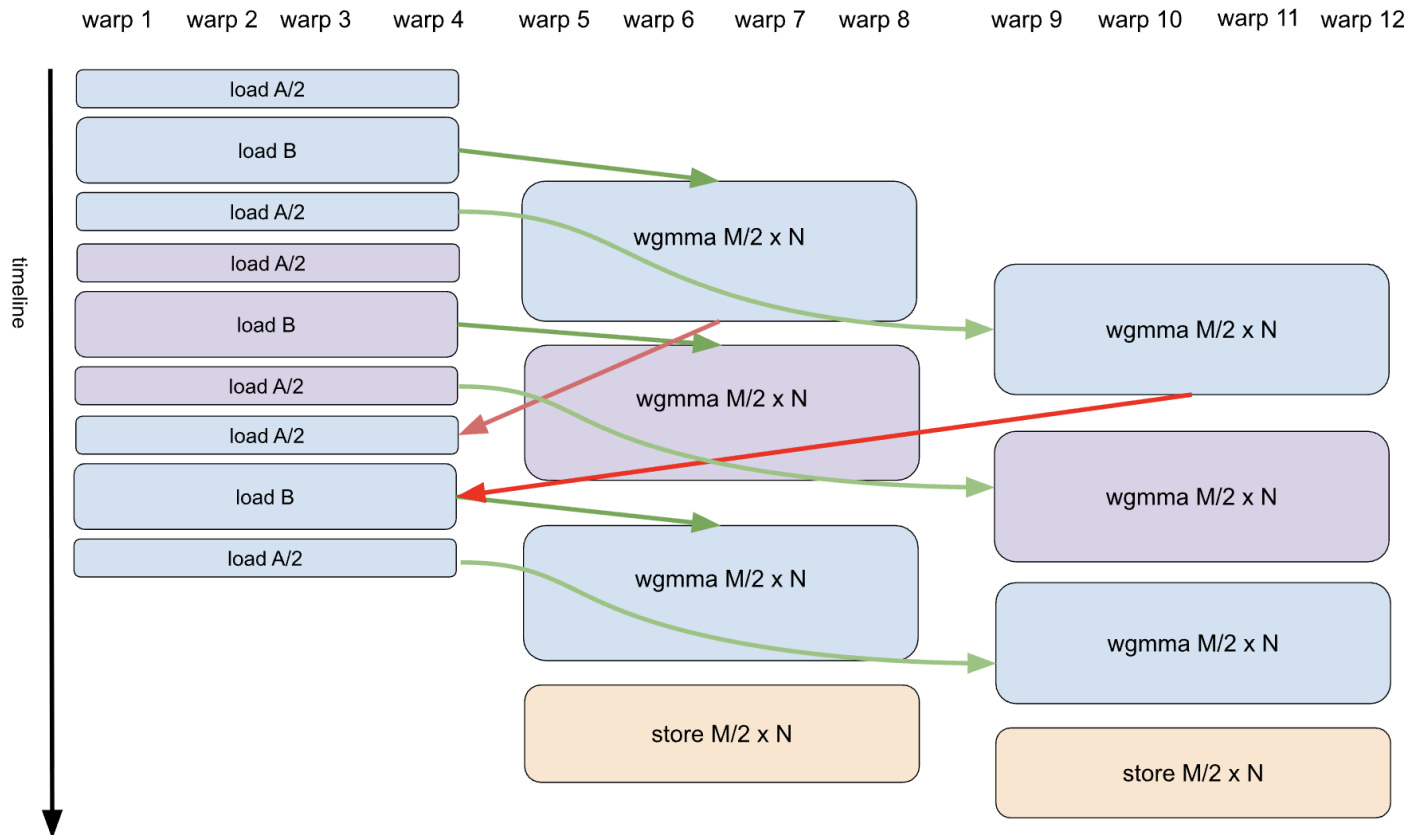
```
# trunk code with num_warps=4
...

with tlx.async_tasks():
    with tlx.async_task("default"):
        ...
    with tlx.async_task(num_warps=1):
        ...

# trunk code
...
```



Cooperative Warp Specialization



```
@triton.jit
def _load(desc, buffer, bar_empty, bar_full, idx, size, offsets, NUM_BUFFERS)
    buf_id = idx % NUM_BUFFERS
    phase = idx // NUM_BUFFERS
    tlx.barrier_wait(bar_empty[buf_id], phase ^ 1)
    tlx.barrier_expect_bytes(bar_full[buf_id], size)
    tlx.async_descriptor_load(desc, x[buf_id], offsets, bar_full[buf_id])
```

```
@triton.jit
def hopper_gemm_ws(a_desc, b_desc, c_desc ... M, N, K, NB: tl.constexpr):

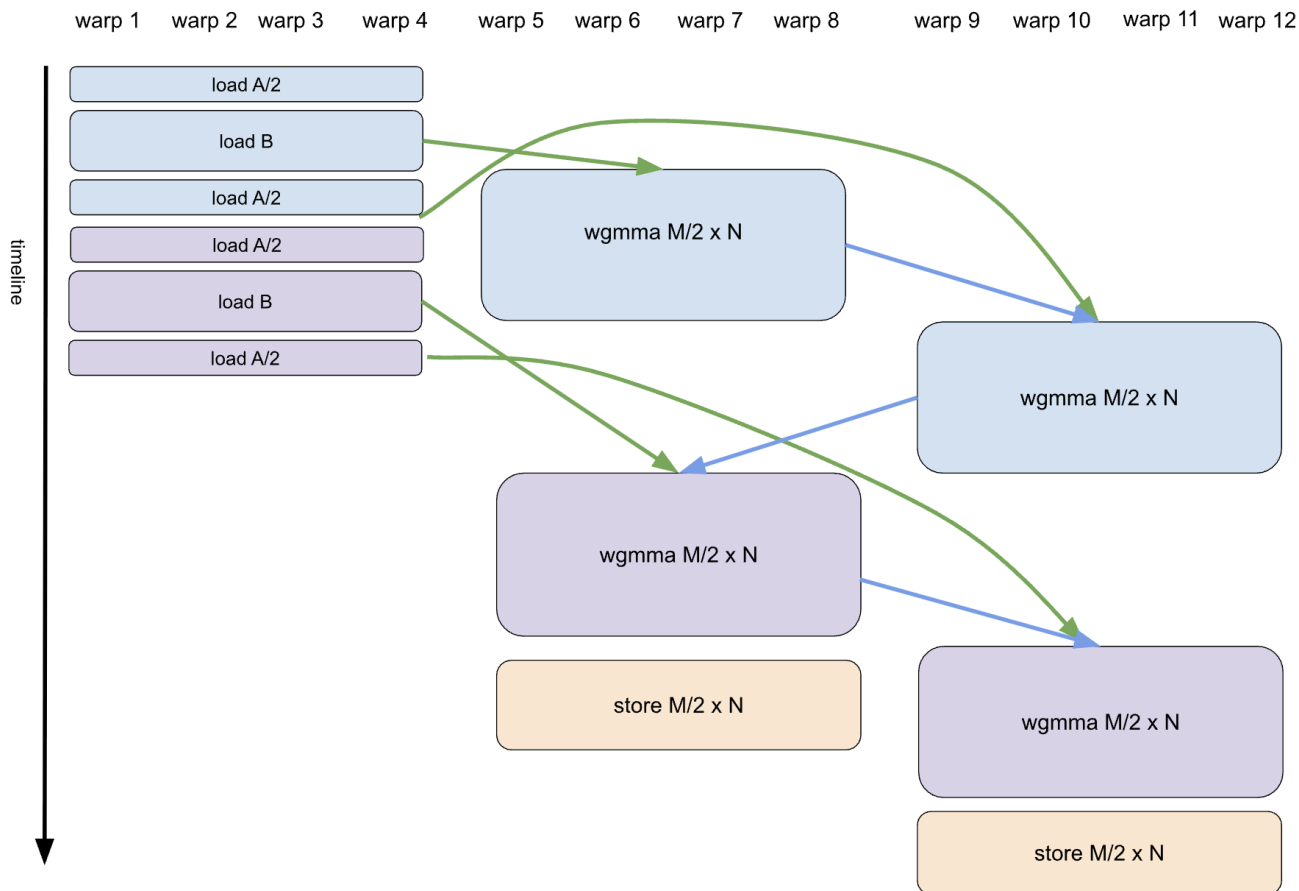
    a = tlx.local_alloc((BLOCK_M // 2, BLOCK_K), tlx.dtype_of(x_desc), NB*2)
    b = tlx.local_alloc((BLOCK_K, BLOCK_N), tlx.dtype_of(y_desc), NB)
    a_full = tlx.alloc_barriers(num_barriers=NB * 2)
    a_empty = tlx.alloc_barriers(num_barriers=NB * 2)
    b_full = tlx.alloc_barriers(num_barriers=NB, arrive_count=2)
    b_empty = tlx.alloc_barriers(num_barriers=NB, arrive_count=2)
```

```
with tlx.async_tasks():
    with tlx.async_task("default"):
        for k in range(0, tl.cdiv(K, BK)):
            _load(a_desc, a, a_empty, a_full, k, BM // 2 * BK, [k * BK])
            _load(b_desc, b, b_empty, b_full, k, BN * BK, [k * BK])
            _load(a_desc, a+NB, a_empty+NB, a_full+NB, k, BM // 2 * BK, [k * BK])
```

```
with tlx.async_task(num_warps=4, num_regs=168, replicate=2):
    cid = tlx.async_task_replica_id()
    for k in range(0, tl.cdiv(K, BK)):
        buf_id = idx % NUM_BUFFERS
        phase = idx // NUM_BUFFERS
        tlx.barrier_wait(a_full[buf_id + cid * NB], phase)
        tlx.barrier_wait(b_full[buf_id], phase)
        acc = tlx.async_dot(a[buf_id], b[buf_id], acc)
        acc = tlx.async_dot_wait(0, acc)
        tlx.barrier_arrive(a_empty[buf_id + cid * NB])
        tlx.barrier_arrive(b_empty[buf_id])
        o_desc.store([offset_am, offset_bn], acc)
```



Ping-Pong Schedule



```
with tlx.async_tasks():
    with tlx.async_task("default") :
        for k in range(0, tl.cdiv(K, BK))
            _load(a_desc, a, a_empty, a_full, k, BM // 2 * BK, [k * BK])
            _load(b_desc, b, b_empty, b_full, k, BN * BK, [k * BK])
            _load(a_desc, a+NB, a_empty+NB, a_full+NB, k, BM // 2 * BK, [k * BK])
```

```
with tlx.async_task(num_warps=4, num_regs=168, replicate=2):
```

```
    cid = tlx.async_task_replica_id()
```

```
    if cid == 1:
```

```
        tlx.named_barrier_arrive(9, 256)
```

```
    for k in range(0, tl.cdiv(K, BK))
```

```
        buf_id = idx % NUM_BUFFERS
```

```
        phase = idx // NUM_BUFFERS
```

```
        tlx.barrier_wait(a_full[buf_id + cid * NB], phase)
```

```
        tlx.barrier_wait(b_full[buf_id], phase)
```

```
    if cid == 0:
```

```
        tlx.named_barrier_wait(9, 256)
```

```
    else:
```

```
        tlx.named_barrier_arrive(10, 256)
```

```
    acc = tlx.async_dot(a[buf_id], b[buf_id], acc)
```

```
    if cid == 0:
```

```
        tlx.named_barrier_arrive(10, 256)
```

```
    else:
```

```
        tlx.named_barrier_arrive(9, 256)
```

```
    acc = tlx.async_dot_wait(0, acc)
```

```
    tlx.barrier_arrive(a_empty[buf_id + cid * NB])
```

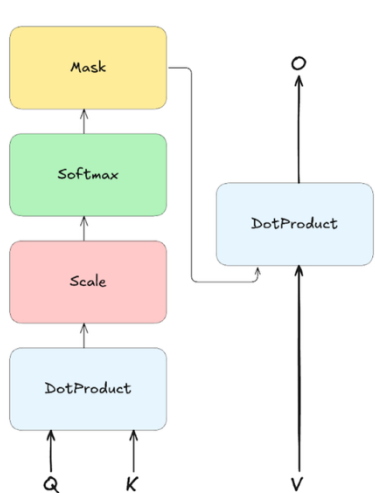
```
    tlx.barrier_arrive(b_empty[buf_id])
```

```
    o_desc.store([offset_am, offset_bn], acc)
```

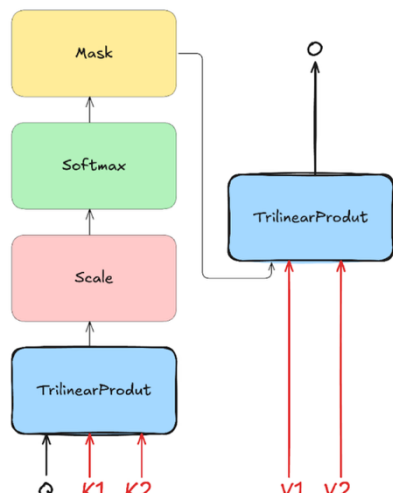


TLX Kernel Performance Study

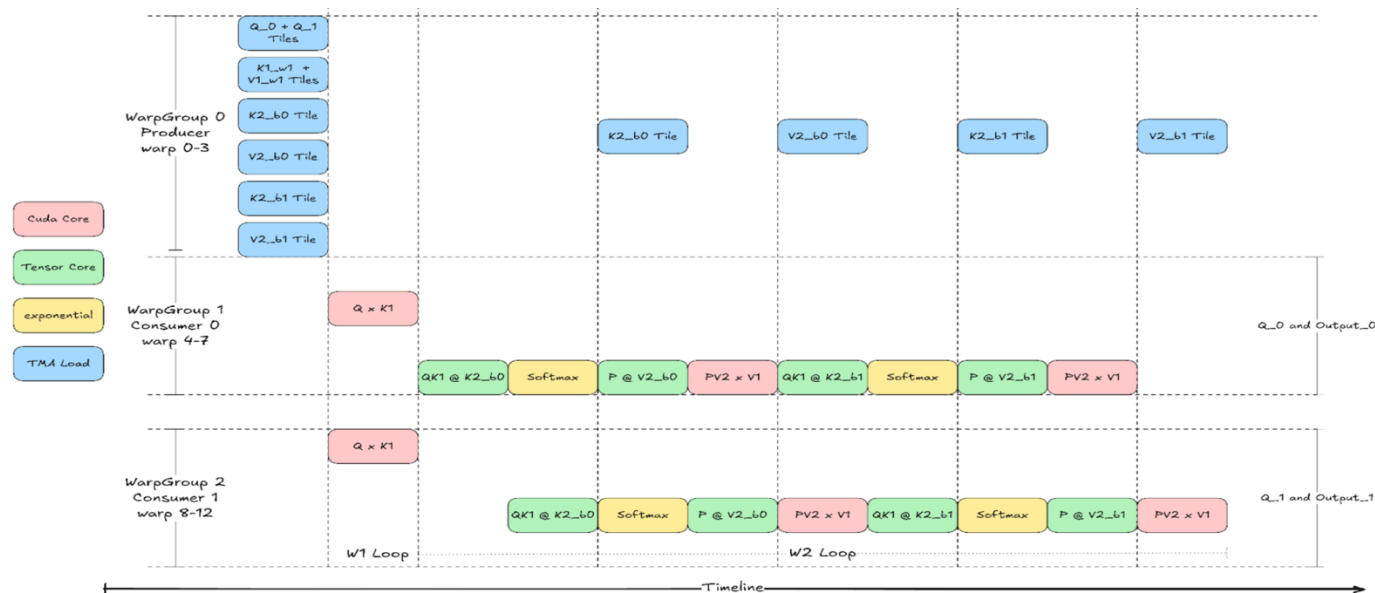
2-Simplicial Attention on H100: <https://pytorch.org/blog/fast-2-simplicial-attention-hardware-efficient-kernels-in-tlx/>



Standard Dot-Product Attention



2-Simplicial Attention



TLX Kernels

	Triton	TLX
2-Simplicial Attention on H100	337 TFLOPS	588 TFLOPS warp specialization computation pipelining ping-pong K1/V1 buffer prefetch
Flash Attention Fwd on H100	446 TFLOPS	548 TFLOPS warp specialization computation pipelining ping-pong
Flash Attention Fwd on B200	789 TFLOPS warp specialization	930 TFLOPS warp specialization mma/load pipelining fine-grain barrier on tile slices



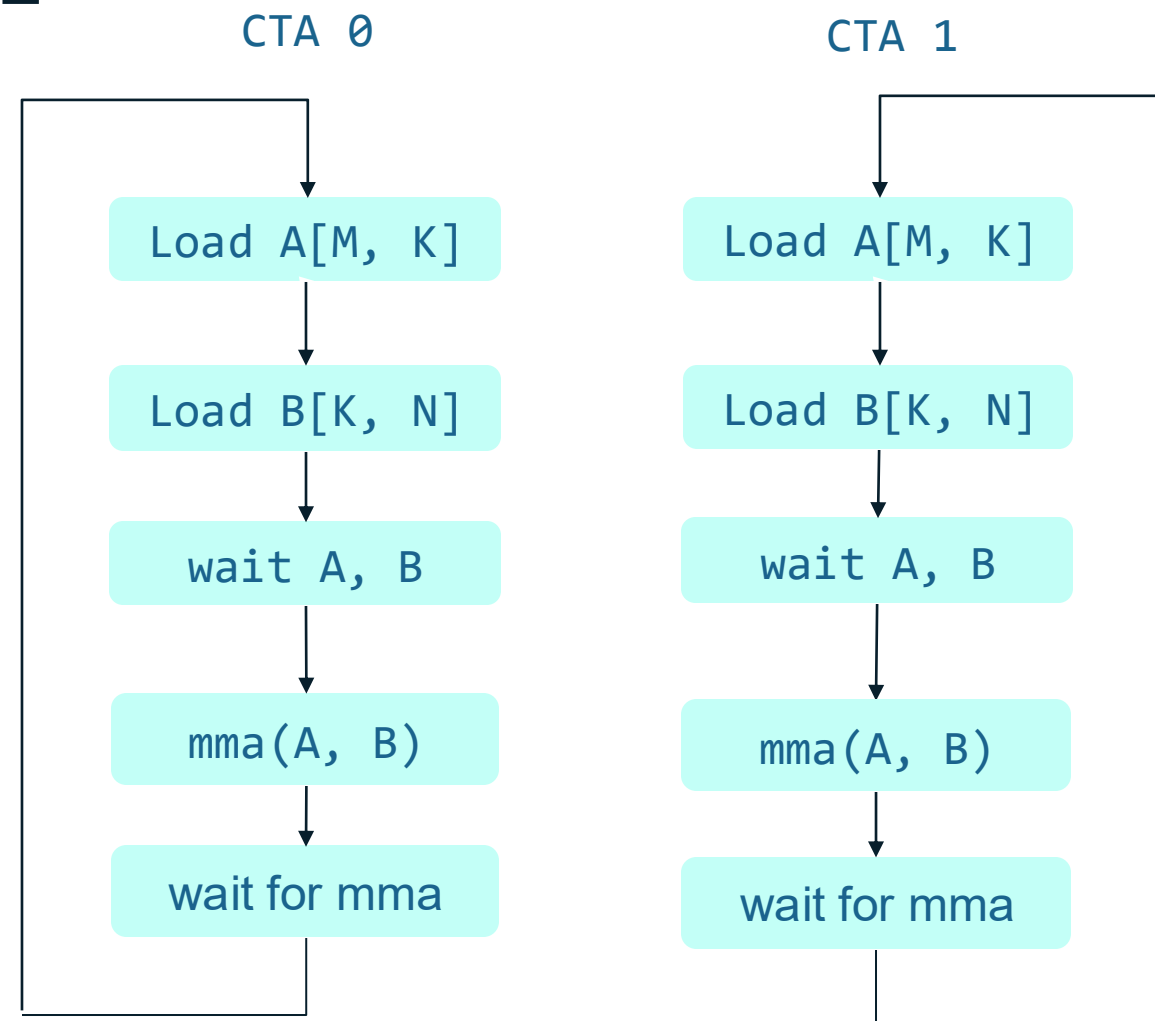
Where We're Going

- More hardware features
 - Paired-CTA
 - Cluster launch control (CLC)
 - Distributed shared memory
- Supporting AMDGPU
 - Same user interface
 - Software-based wave specialization support



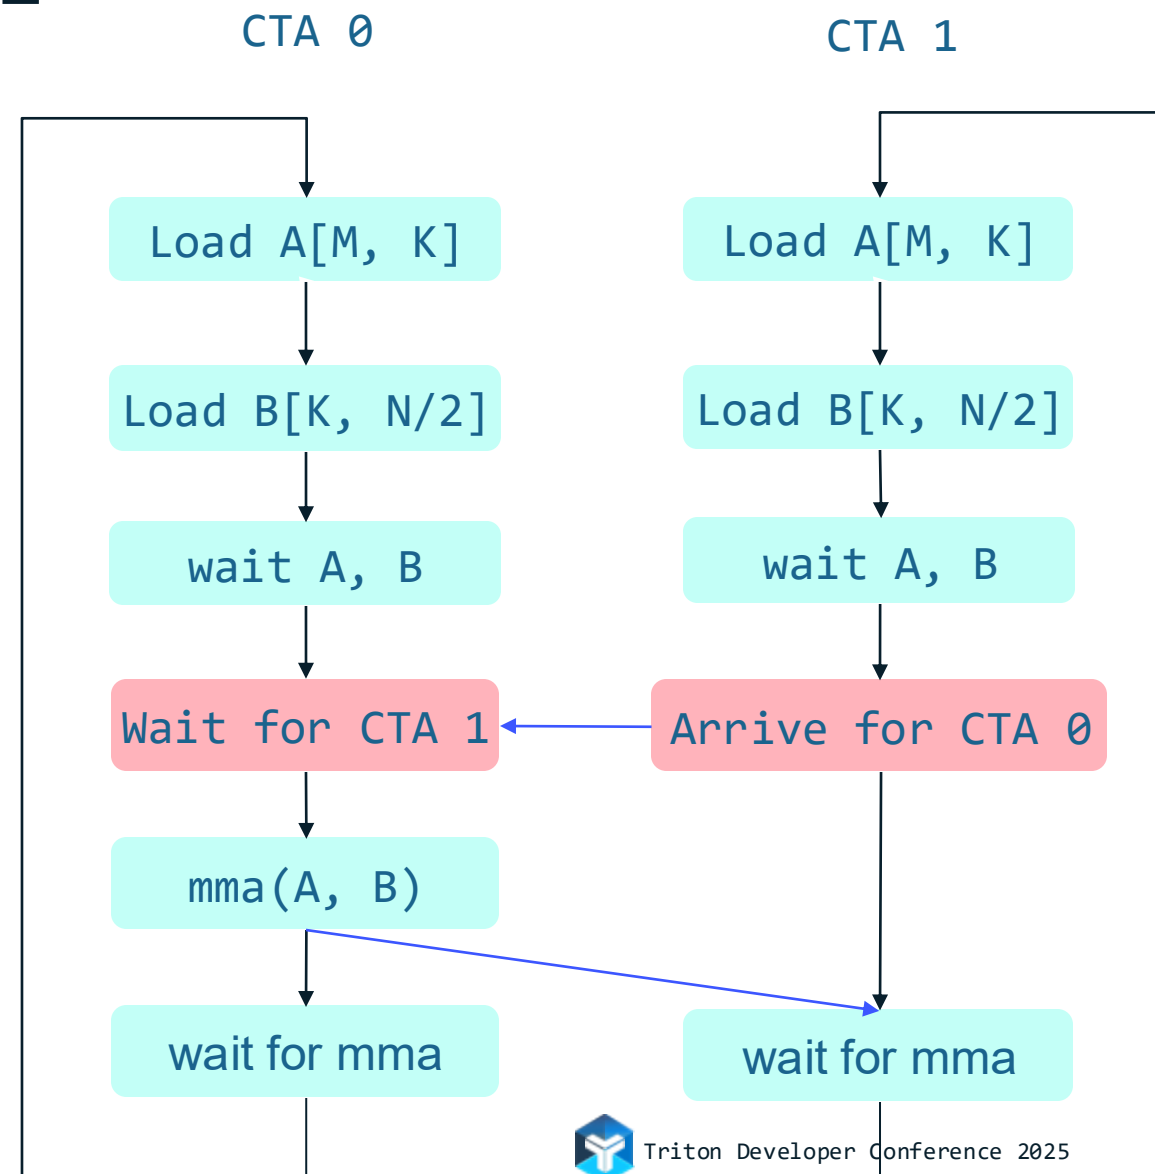
Paired-CTA MMA on Blackwell

- Single-CTA
 - Each CTA computes $M*N$
 - Two CTAs compute $2*M*N$



Paired-CTA MMA on Blackwell

- Paired-CTA
 - Each CTA computes $M*N$
 - Two CTAs co-compute $2*M*N$
 - B is shared, half stored in CTA 0, half in CTA 1
 - Only one extra barrier arrive/wait
- Less memory traffic
- Less shared memory usage
- +2% speedup over single-CTA
- 98% of cuBLAS GEMM performance



Cluster Launch Control (CLC)

- Enables dynamic tile scheduling for imbalanced workloads



```
start_pid = tl.program_id(axis=0)
phase = 0
while start_pid != -1 :
    # compute tile offset based on start_pid
    ...
    # GEMM K-loop
    for k in range(0, tl.cdiv(K, BLOCK_SIZE_K)):
        ...

    # Issue async clc.try_cancel for the next
    # available CTA
    tlx.barrier_expect_bytes(clc_mbars, 16)
    tlx.clc_issue(clc_responses, clc_mbars)
    tlx.barrier_wait(clc_mbars, phase)


    # Extract CTA ID from CLC response
    start_pid = tlx.clc_query(clc_responses)
    phase = phase ^ 1
```

dynamic launched kernel:


Step 1: Each available SM launches its first working tile with matched CTA ID *Kernel launch*



Step 2: If one SM (SM 0 for example) finishes its working tile, it can fetch the next available CTA ID with CLC without launching a new kernel



Fetch available CTA ID



Main Takeaways

- TLX = Triton Low-level Language Extension
 - Hardware-level control within the Triton ecosystem
- Available now!
 - Download : <https://github.com/facebookexperimental/triton> or QR code below.
 - Spec: <https://github.com/facebookexperimental/triton/blob/main/README.md>
 - Barrier manual:
https://github.com/facebookexperimental/triton/blob/main/third_party/tlx/tlx_barriers.md
- Supports NVIDIA Hopper/Blackwell
 - AMD GPU support is in development.
- Try it out and tell us what you think!





Triton Developer Conference 2025

Hongtao Yu (Meta)

